

# HW#3

Halid Kopanski

2022-06-06

## LMR 4.1

```
##
## Call:
## lm(formula = lpsa ~ ., data = prostate)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7331 -0.3713 -0.0170  0.4141  1.6381
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.669337   1.296387   0.516  0.60693
## lcavol      0.587022   0.087920   6.677 2.11e-09 ***
## lweight     0.454467   0.170012   2.673  0.00896 **
## age        -0.019637   0.011173  -1.758  0.08229 .
## lbph       0.107054   0.058449   1.832  0.07040 .
## svi        0.766157   0.244309   3.136  0.00233 **
## lcp       -0.105474   0.091013  -1.159  0.24964
## gleason     0.045142   0.157465   0.287  0.77503
## pgg45       0.004525   0.004421   1.024  0.30886
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7084 on 88 degrees of freedom
## Multiple R-squared:  0.6548, Adjusted R-squared:  0.6234
## F-statistic: 20.86 on 8 and 88 DF,  p-value: < 2.2e-16
```

a)

```
##          fit          lwr          upr
## 1 2.389053 2.172437 2.605669

##          fit          lwr          upr
## 1 2.389053 0.9646584 3.813447

## [1] 63.86598
```

b)

```
##          fit          lwr          upr
## 1 3.17454 2.270398 4.078682
```

```
##          fit          lwr          upr
## 1 3.17454 1.501384 4.847695
```

c)

```
##
## Call:
## lm(formula = lpsa ~ lcavol + lweight + svi, data = prostate)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.72964 -0.45764  0.02812  0.46403  1.57013
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -0.26809    0.54350  -0.493  0.62298
## lcavol         0.55164    0.07467   7.388 6.3e-11 ***
## lweight        0.50854    0.15017   3.386 0.00104 **
## svi            0.66616    0.20978   3.176 0.00203 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7168 on 93 degrees of freedom
## Multiple R-squared:  0.6264, Adjusted R-squared:  0.6144
## F-statistic: 51.99 on 3 and 93 DF,  p-value: < 2.2e-16
```

```
##          fit          lwr          upr
## 1 2.372534 2.197274 2.547794
```


```
##          fit          lwr          upr
## 1 2.372534 0.9383436 3.806724
```

```
## Analysis of Variance Table
##
## Model 1: lpsa ~ lcavol + lweight + svi
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + gleason +
##          pgg45
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      93 47.785
## 2      88 44.163   5    3.6218 1.4434 0.2167
```

## LMR 6.1

```
##
## Call:
## lm(formula = total ~ expend + salary + ratio + takers, data = sat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -90.531 -20.855  -1.746  15.979  66.571
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1045.9715    52.8698   19.784 < 2e-16 ***
## expend         4.4626    10.5465    0.423  0.674
## salary        1.6379     2.3872    0.686  0.496
## ratio        -3.6242     3.2154   -1.127  0.266
## takers       -2.9045     0.2313  -12.559 2.61e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 32.7 on 45 degrees of freedom
## Multiple R-squared:  0.8246, Adjusted R-squared:  0.809
## F-statistic: 52.88 on 4 and 45 DF,  p-value: < 2.2e-16
```

a)



Results of the Breusch - Pagan Test


```
##
```

```
## Breusch-Pagan test
##
## data: sat_model1
## BP = 2.7234, df = 4, p-value = 0.6051
```

The results of the Non Constant Variance Test

```
## Non-constant Variance Score Test
## Variance formula: ~ fitted.values
## Chisquare = 0.6972119, Df = 1, p = 0.40372
```

b)






Results of the Shapiro and the Durbin-Watson Tests

```
##
## Shapiro-Wilk normality test
##
## data: sat_model1$residuals
## W = 0.97691, p-value = 0.4304
```

```
##
## Durbin-Watson test
##
## data: sat_model1
## DW = 2.4525, p-value = 0.9459
## alternative hypothesis: true autocorrelation is greater than 0
```


c)





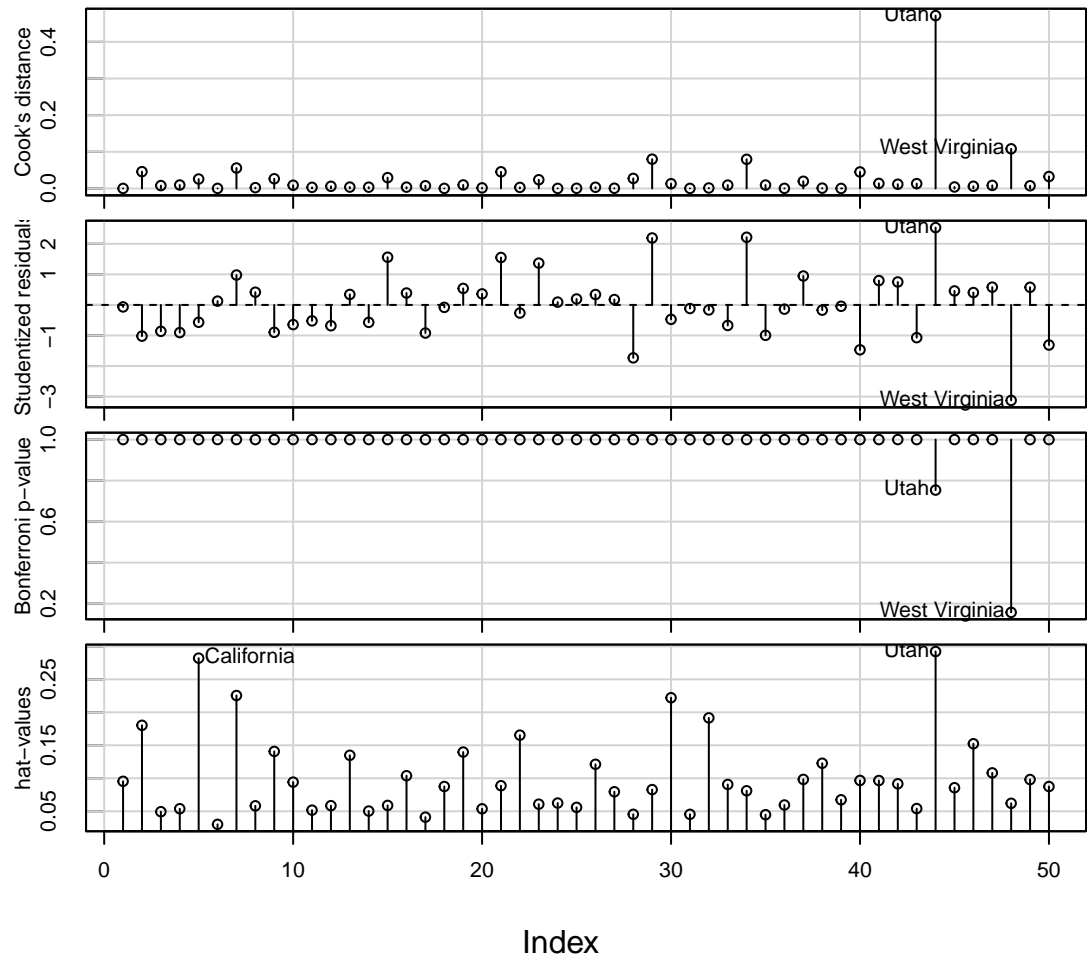
d)

```
## No Studentized residuals with Bonferroni p < 0.05
## Largest |rstudent|:
##          rstudent unadjusted p-value Bonferroni p
## West Virginia -3.124428      0.0031496      0.15748
```




e)

Diagnostic Plots







## LMR 6.8

a)

```
## [1] 1.00000 17.47144 25.30482 58.60610 83.59121 100.63222 137.89717
## [8] 175.28623 192.61449 213.00748 228.15747 268.20620 555.67072
```

```
##                                vif_x
## [1,] "age"                    "2.25045023586224"
## [2,] "weight"                 "33.50931979125"
## [3,] "height"                 "1.67459083959811"
## [4,] "neck"                   "4.32446326424943"
## [5,] "chest"                  "9.46087732137424"
## [6,] "abdom"                  "11.7670733753212"
## [7,] "hip"                    "14.796519836923"
## [8,] "thigh"                  "7.77786469219344"
## [9,] "knee"                   "4.61214673600487"
## [10,] "ankle"                 "1.90796099507021"
## [11,] "biceps"                "3.61974357536549"
## [12,] "forearm"               "2.1924921166644"
## [13,] "wrist"                 "3.37751489619545"
```

b)

```
## [1] 1.00000 18.39787 26.21547 61.53224 91.07633 114.44792 148.72518
## [8] 178.80871 202.08708 211.78359 240.69468 276.35018 554.79777
```

```
##                                vif_x2
## [1,] "age"                    "2.27819073931863"
## [2,] "weight"                 "45.2988434683429"
## [3,] "height"                 "3.43958730850223"
## [4,] "neck"                   "3.97889794019821"
## [5,] "chest"                  "10.7125052438853"
## [6,] "abdom"                  "11.9675796797964"
## [7,] "hip"                    "12.1462491415389"
## [8,] "thigh"                  "7.15371104845086"
## [9,] "knee"                   "4.44175243874575"
## [10,] "ankle"                 "1.81025271253279"
## [11,] "biceps"                "3.40952410870031"
## [12,] "forearm"               "2.42287790935087"
## [13,] "wrist"                 "3.26367710797482"
```

c)

```
##
## Call:
## lm(formula = brozek ~ age + weight + height, data = new_fat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.0260  -3.6537   0.0569   3.7588  11.9011
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 54.31985    9.63347   5.639 4.69e-08 ***
## age         0.12575    0.02599   4.838 2.31e-06 ***
## weight      0.23519    0.01373  17.124 < 2e-16 ***
## height     -1.18089    0.14638  -8.067 3.17e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.986 on 246 degrees of freedom
## Multiple R-squared:  0.5838, Adjusted R-squared:  0.5787
## F-statistic: 115 on 3 and 246 DF, p-value: < 2.2e-16

## [1] 1.00000 13.87911 25.03771

##           vif_x3
## [1,] "age"      "1.08330491018921"
## [2,] "weight"   "1.3811645842945"
## [3,] "height"   "1.46966479387396"
```

d)

```
##           fit      lwr      upr
## 1 18.48834 8.647863 28.32882
```

e)

```
##           fit      lwr      upr
## 1 20.18367 10.32046 30.04688
```

f)

```
##           fit      lwr      upr
## 1 3.720148 -6.28208 13.72238
```

## Code

```
library(faraway)
library(car)
library(MASS)
library(lmtest)
library(tidyverse)
pros_model1 <- lm(lpsa ~ ., data = prostate)
summary(pros_model1)
new_patient <- tibble(lcavol = 1.44692, lweight = 3.62301, age = 65.00000,
                      lbph = 0.30010, svi = 0.00000, lcp = -0.79851,
                      gleason = 7.0000, pgg45 = 15.0000)

predict(pros_model1, newdata = new_patient, interval = "confidence")
predict(pros_model1, newdata = new_patient, interval = "prediction")
mean(prostate$age)
new_patient1 <- tibble(lcavol = 1.44692, lweight = 3.62301, age = 25.00000,
                      lbph = 0.30010, svi = 0.00000, lcp = -0.79851,
                      gleason = 7.0000, pgg45 = 15.0000)

predict(pros_model1, newdata = new_patient1, interval = "confidence")
predict(pros_model1, newdata = new_patient1, interval = "prediction")
pros_model2 <- lm(lpsa ~ lcavol + lweight + svi, data = prostate)
summary(pros_model2)
predict(pros_model2, newdata = new_patient1, interval = "confidence")
predict(pros_model2, newdata = new_patient1, interval = "prediction")
print(anova(pros_model2, pros_model1))
sat_model1 <- lm(total ~ expend + salary + ratio + takers, data = sat)
summary(sat_model1)
plot(x = sat_model1$fitted.values, y = sat_model1$residuals, pch = 19,
     main = "Residuals vs Fitted Values", xlab = "Fitted Values", ylab = "Residuals")
res_line <- loess(sat_model1$residuals ~ sat_model1$fitted.values)
j <- order(sat_model1$fitted.values)
lines(sat_model1$fitted.values[j], res_line$fitted[j], col = "red", lwd = 3)
abline(h = 0)
bptest(sat_model1, studentize = FALSE)
ncvTest(sat_model1)
qqnorm(sat_model1$residuals, ylab = "Residuals", main = "Q-Q Plot (Normality)", pch = 19)
qqline(sat_model1$residuals)
shapiro.test(sat_model1$residuals)
dwtest(sat_model1)
X <- sat %>% select(expend, salary, ratio, takers)
X <- cbind(rep(1, nrow(X)), X)

X <- data.matrix(X)

#print(X)
Xt <- t(X)

XtX_inv <- solve(Xt %*% X)

XtY <- Xt %*% sat$total
```

```

beta_hat <- XtX_inv %*% XtY

P <- X %*% XtX_inv %*% Xt

res_mean <- mean(sat_model1$residuals)
res_sd <- sd(sat_model1$residuals)
stan_res <- (sat_model1$residuals - res_mean) / res_sd

plot(x = diag(P),
     y = sat_model1$residuals,
     pch = 19,
     main = "Residuals vs Leverage",
     xlab = "Leverage",
     ylab = "Residuals",
     xlim = c(0, 0.35))
text(diag(P), sat_model1$residuals, rownames(P),
     cex = 0.6, pos = 4, col = "red")
hat_data <- cbind(1:50, diag(P))

plot(hat_data, pch = 19, xlab = "", ylab = "hii values")
abline(h = 2 * 5 / 50)
text(hat_data[,1], hat_data[,2], names(diag(P)),
     cex = 0.5, pos = 4, col = "red")
dotchart(sort(diag(P)), pch = 19, cex = 0.5)
outlierTest(sat_model1)

plot(x = sat_model1$fitted.values, y = rstandard(sat_model1), pch = 19,
     main = "Standardized Residuals vs Fitted Values", xlab = "Fitted Values", ylab = "Standardized Residuals")
abline(h = c(0, 2, -2), col = "blue", lty = 4, lwd = 2)
text(x = sat_model1$fitted.values, y = rstandard(sat_model1), rownames(sat),
     cex = 0.5, pos = 4, col = "red")
#cooks.distance(sat_model1)

halfnorm(diag(P), labs = rownames(P), ylab = "Leverages")
abline(h = 2 * sat_model1$rank / nrow(sat))

SST <- sum((sat$total - mean(sat$total))^2)
SSR <- sum((sat_model1$fitted.values - mean(sat$total))^2)
SSE <- sum((sat$total - sat_model1$fitted.values)^2)
MSR <- SSR/(sat_model1$rank - 1)
MSE <- SSE/(nrow(sat) - sat_model1$rank)
influenceIndexPlot(sat_model1)
cooks_dis_calc <- sat_model1$residuals^2 /
  (sat_model1$rank * MSE * (diag(P)/(1-diag(P))^2))

dotchart(sort(cooks_dis_calc), cex = 0.5, pch = 19)
fat_model1 <- lm(brozek ~ age + weight + height + neck + chest +
  abdom + hip + thigh + knee + ankle + biceps +
  forearm + wrist, data = fat)

X <- model.matrix(fat_model1)[,-1]
eigen_X <- eigen(t(X) %*% X)

```

```

sqrt(eigen_X$values[1]/eigen_X$values)

vif_x <- rep(0, ncol(X))

for(i in 1:ncol(X)){
  vif_x[i] <- 1 / (1 - summary(lm(X[,i] ~ X[,-i]))$r.squared)
}

vif_x <- cbind(colnames(X), vif_x)
print(vif_x)
new_fat <- fat[-c(39,42),]

fat_model2 <- lm(brozek ~ age + weight + height + neck + chest +
  abdom + hip + thigh + knee + ankle + biceps +
  forearm + wrist, data = new_fat)

X2 <- model.matrix(fat_model2)[,-1]
eigen_X2 <- eigen(t(X2) %*% X2)

sqrt(eigen_X2$values[1]/eigen_X2$values)

vif_x2 <- rep(0, ncol(X2))

for(i in 1:ncol(X2)){
  vif_x2[i] <- 1 / (1 - summary(lm(X2[,i] ~ X2[,-i]))$r.squared)
}

vif_x2 <- cbind(colnames(X2), vif_x2)
print(vif_x2)

fat_model3 <- lm(brozek ~ age + weight + height, data = new_fat)
summary(fat_model3)

X3 <- model.matrix(fat_model3)[,-1]
eigen_X3 <- eigen(t(X3) %*% X3)

sqrt(eigen_X3$values[1]/eigen_X3$values)

vif_x3 <- rep(0, ncol(X3))

for(i in 1:ncol(X3)){
  vif_x3[i] <- 1 / (1 - summary(lm(X3[,i] ~ X3[,-i]))$r.squared)
}

vif_x3 <- cbind(colnames(X3), vif_x3)
print(vif_x3)

fat_1 <- as.data.frame(t(apply(X3, 2, median)))

```

```
predict(fat_model3, newdata = fat_1, interval = "prediction")

fat_2 <- as.data.frame(cbind(age = 40, weight = 200, height = 73))
predict(fat_model3, newdata = fat_2, interval = "prediction")

fat_3 <- as.data.frame(cbind(age = 40, weight = 130, height = 73))
predict(fat_model3, newdata = fat_3, interval = "prediction")
```